THE WHAT, BEFORE THE HOW

Let me show you what it does

A $4.8B gas processing expansion at Jubail. 400,000+ documents across Aconex, P6 and daily reports. One agent, reading continuously.

1 · DETECTION

The agent finds a $42.3M exposure nobody logged — 6 days before the notice bar

2 · EVIDENCE CHAIN

Six documents, linked from the contract clause to the daily report to the schedule diff

3 · DRAFT NOTICE

Every sentence sourced. A named human approves — the agent never sends

4 · ORCHESTRATION

Approval triggers ComplyNova + FlowNova + Schedule Interrogator in NovaStudio

▶ **Open live prototype**

THE PUNCHLINE

From reconstructing the past → to reading the record as it happens

TODAY	WITH ASSET NOVA
Contracts team manually reconstructs 18 months of correspondence to assess one claim	Agent reads 400K documents continuously, surfaces events as they happen
Notice deadlines missed because nobody connected the letter to the clause to the schedule	Deadlines computed from the contract; countdown visible; draft notice ready
$60M+ disputes that take a year to resolve	Disputes prevented by proper contract administration — the #1 cause for a decade (Arcadis)
Monthly reports assembled by copy-paste across 6 systems, 6 weeks stale	Live, evidence-linked project intelligence, rendered on demand
One expert planner reads two XER files for a week to find hidden changes	Agent forensically diffs 412 logic changes and cross-examines them in minutes

WHY THIS MATTERS

Capital megaprojects fail at industrial scale

89%

of Middle East O&G projects finish over budget (EY, 365 projects)

+59%

average O&G megaproject cost vs approved FID budget

33%

of project capex ends up disputed — $95B across 2,200+ projects (HKA CRUX)

$91M

average Middle East dispute; 12.5 months to resolve (Arcadis)

> Over 75% of respondents said the PM or Engineer's conduct was at the heart of how disputes crystallized — most often **"a lack of understanding of the procedural aspects of the contract."** Over 60% said **proper contract administration** would have the single largest impact in avoiding their disputes.
>
> — Arcadis, Global Construction Disputes Report

The data already exists — in letters, minutes, schedules. No human team can read 400,000 documents and connect every clause to every event to every deadline. **This is a reading problem, not a judgment problem.**

THE STRUCTURAL INSIGHT

Software digitized the silos. It never connected them.

```
SCHEDULE (P6)          COST (SAP)          DOCUMENTS (Aconex)
    │                      │                        │
    └────── each contractor runs its own instance ──┘
                           │
                 ┌─────────┴──────────────┐
                 │   THE REAL INTEGRATION  │
                 │   LAYER OF THE INDUSTRY: │
                 │   EXCEL + POWERPOINT     │
                 │   assembled monthly      │
                 └──────────────────────────┘
```

The failure mode is **inter-party, cross-silo, and unstructured**. Previous software couldn't read text at scale. LLM-based agents can. That's the structural unlock — and it's why this problem is solvable now and wasn't in 2015.

COMPETITIVE LANDSCAPE

Every AI player misses the real problem

PLAYER	WHAT THEY DO	WHAT THEY DON'T DO
ALICE Technologies	Schedule simulation + optimization	Read correspondence, handle claims, cross-system
nPlan	ML schedule risk forecasting	Anything beyond the P6 file
Document Crunch (Trimble)	Contract clause extraction	The live project record; static, not continuous
Nodes & Links	Schedule assurance + auditable AI Q&A	Contract mechanics, notices, cost; planner persona only
SuperHive (AI71 / ATRC)	Sovereign-AI studios: design validation, permits, site monitoring	Contract, correspondence, claims. Buildings ICP, not process-industry owners

The open seat: owner PMO, cross-system, full lifecycle, process-industry megaprojects. SuperHive is the most credible future competitor (schedule intelligence in development) — which is exactly why speed on the commercial wedge matters.

PRODUCT THESIS

Not a new product — the capital-projects layer of NovaOS

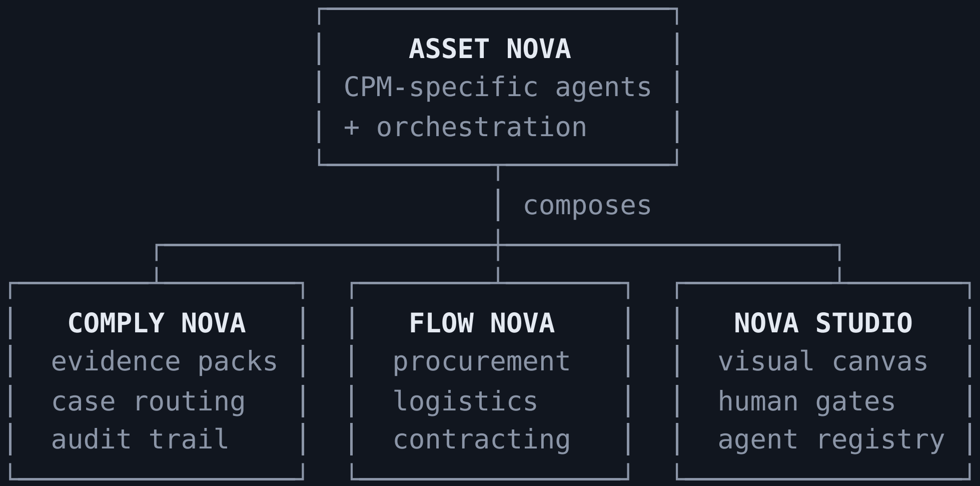

The real product is the **substrate**: a cross-system, document-grounded project record. Agents are composable skills on top. "Skills compound, dashboards decay" — Asset Nova is the purest expression of it: the skill is reading the record; views are rendered on demand and thrown away.

THE WEDGE

Commercial Guardian first — by the value equation

Value = (Dream outcome × Likelihood) ÷ (Time to value × Effort)

	COMMERCIAL GUARDIAN	BUDGET ESTIMATION	SCHEDULE RISK
Dream outcome	Protect $30–80M in disputes/leakage per project	Accurate FID budget	Earlier warning
Likelihood	High — every finding has a document evidence chain	Low — needs outcome data that doesn't exist	Medium
Time to value	Days (runs on historical archives)	Months (next gate)	Monthly (next XER)
Effort to start	Low — document upload, no ERP integration	High (SAP + history DB)	Medium
Verdict	**Start here**	Horizon 3 (crown jewel)	Fast-follow

HOW IT WORKS

Agent architecture — deterministic where it must be

```
INPUTS (read-only)                AGENT PIPELINE                    OUTPUTS (human-approved)
__________________                ______________                    ________________________

Contracts + amendments ───────> OBLIGATION EXTRACTOR ──────────> clause registry + deadline
Aconex/CDE correspondence          (contract form parser)            rules (DETERMINISTIC)
(letters, RFIs, minutes)                   |
P6 XER files ──────────────────> EVENT DETECTOR ─────────────────> events + evidence chains
Daily / weekly reports             (cross-system correlation,        → weekly review ritual
                                    LLM reasoning)
                                           |
                                 NOTICE DRAFTER ─────────────────> draft notices w/ citations
                                   (generation + citation check)     → named human approves → sent
```

MODEL

Open-weights (Llama/Qwen-class), fine-tuned on construction correspondence. Weights, not wrappers.

DEPLOYMENT

Sovereign GPU, in-Kingdom VPC. SACS-002 compliant, PDPL-aligned. Data never leaves.

AUDIT

Full log of every read and inference. Glass Box by design.

SUCCESS METRICS

Three layers — business, product, AI engineering

BUSINESS — DOES IT PAY?

- $20M+ validated exposures surfaced per pilot (customer countersigned)
- Zero notice deadlines missed on covered projects
- $10–12M ARR at 24 months (5 pilots → 3 conversions → 6 projects)

PRODUCT — IS IT TRUSTED?

- 90%+ weekly review-ritual adoption — the habit IS the product
- 80%+ surfaced events actioned, not just seen
- 70%+ draft notices approved with minor/no edits

AI ENGINEERING — IS IT RELIABLE?

- 85%+ event-detection precision (false alarms kill trust)
- 99.5%+ evidence-chain accuracy — deterministic citation verifier
- Zero hallucinations in sent notices (entailment check + 100% human review)
- Deadline math = rules engine, not LLM. <$0.05/doc blended cost

Kill criteria: precision < 60% after 8 weeks, or the team stops reviewing for 4 straight weeks → pause, return to backtest.

GOVERNANCE

The three questions, answered before deployment

QUESTION	ANSWER FOR COMMERCIAL GUARDIAN
Who owns the agent's decisions?	Named Contracts Lead per project. Agent owner = PMO Director. Orbitron owns model behavior within contract.
What does it do outside its scope?	Nothing. Refuses, logs, escalates. Below-confidence items queue for triage. Escalation: agent → Contracts Lead → PMO → steering committee.
How do we measure value, not activity?	Dollar-validated exposures (customer countersigned). Notices sent on time. Edit-distance on drafts. **Documents processed is NOT a KPI.**

Agentic AI is organizational redesign, not a technical rollout. Governance isn't the appendix — it's the product.

ROADMAP

Three horizons, one compounding substrate

H1 · 0-6 MO — "FIND THE MONEY LEAKS"

- **Commercial Guardian** — claims + notices
- **Schedule Interrogator** — XER forensics, early warning
- Substrate v1: contracts + correspondence indexed & linked

H2 · 6-18 MO — "RUN THE MONTHLY GRIND"

- **EAC Reconciler** — live cost forecast, drafts monthly report
- **Deliverable Chaser** — critical-path-ranked chasing
- Substrate v2: + cost + document register (daily habit = can't leave)

H3 · 18+ MO — "GATEKEEPER OF CAPITAL"

- **FID Red Team** — audits new proposals vs portfolio history
- **Handover Assembler** — bridge to O&M / digital twin
- Substrate v3: portfolio outcomes — the customer's own alpha, kept

Each finished project's data trains the next FID review. **That's the flywheel.**

GO TO MARKET

Land a project team, not a procurement cycle

OLD ENTERPRISE AI SALES	ASSET NOVA GTM
18-month RFP cycle	Phase 0 backtest on a closed project (6 wks, no live risk) → paid pilot on 1–2 active projects
Per-seat SaaS pricing	Outcome-anchored: ~0.1–0.3% of project value. On a $5B program = $5–15M. "ROI guaranteed" motion
Cloud-first, residency afterthought	In-Kingdom VPC / on-prem day 1. SACS-002. IKTVA-ready. Non-negotiable for Aramco-form contracts
Sell to IT	Sell to the Capital Projects Director who answers to the board for FID promises

One honest trade-off: we don't replace P6, SAP or Aconex — we read them. If the source record is garbage, we read garbage faster. That's fine: on megaprojects the record is rich. It's just unread. Land: one project team · Expand: program EPC packages · Defend: the substrate moat

RISKS — STRESS-TESTED BY 5 REVIEWER AGENTS

What could kill this, and what I'd do about it

RISK	SEV	MITIGATION
Exposure register is **discoverable in arbitration** — "known but unactioned" is ammunition	Critical	Customer counsel as agent owner; customer-controlled retention; team dispositions; privilege labeling
Deadline computed wrong, entitlement lost	Critical	Deterministic rules engine; mandatory customer verification; liability cap + E&O; decision-support disclaimer
Value math stacks assumptions — **"surfaced" ≠ "recovered"**	Serious	Anchor on backtest: one real lapsed time-bar with known arbitration outcome
Contracts teams **don't want their work graded**	Serious	Frame as workload relief; backtest closed projects first; weekly triage cap
Unauthorized practice of law risk in KSA	Serious	Document automation under counsel supervision; authorized-signatory approval; local counsel opinion

EXECUTION

My first 90 days as PM

1-30 · VALIDATE

- 5–8 discovery interviews: owner PMO leads + contracts managers in KSA
- Competitor watch cadence (SuperHive schedule studio)
- v0.1 scope: one contract form (FIDIC Yellow + Aramco amendments), one source (Aconex export)

31-60 · BUILD & BACKTEST

- Ship v0.1 on one closed project
- Measure recall/precision vs expert retrospective audit
- Produce the "here's what you missed" sales artifact

61-90 · PILOT & LEARN

- Live pilot on one active project; embed the weekly review ritual
- Track actioned rate, draft approval rate, triage volume
- First customer-countersigned value memo → program license case

DECISION TREE IF THINGS GO WRONG

Owners won't share archives → pivot to contractor-side wedge · Precision < 60% → narrow to time-bar notices only, retrain · Arabic blocks ingestion → scope v0.1 to English correspondence (~60–70% of the formal record), build Arabic pipeline in parallel

CLOSING

Why me

- [Your specific PM background — fill from personal context]
- I understand the NovaOS architecture — this composes existing agents, it's not a point-solution pitch
- I stress-tested it from five perspectives (engineering, legal, customer, executive, skeptic) before presenting
- The wedge builds the data substrate for a platform, not a feature that dead-ends
- I can ship this: the 90-day plan delivers a customer-validated artifact, not a pitch deck

A1 · IMPACT SIZING

The value math, floor to ceiling

$50M
expected value protected per project per year

2–5×
ROI in the worst case (avoided legal spend on one dispute)

$100M+
best case — one preserved entitlement (Kashagan-class save)

Driver tree: $5B capex → ~$1B addressable disputed value (conservative haircut from HKA's 33%) → 70% detection vs ~40% manual → 3–8% value capture. Full analysis:

`outputs/analyses/impact-sizing-commercial-guardian-2026-08-26.md`

ARR bridge (bottoms-up): 40 target accounts → 12 qualified → 5 paid pilots ($300–500K) → 3 program licenses → 6 projects live at $1.5–3M ACV = **$10–12M ARR at month 24**. One Aramco program-wide deal (Jafurah-scale, 10+ EPC packages) changes the curve entirely.

A2 · SUPERHIVE (AI71) DEEP-DIVE

The most credible future competitor

WHO

Enterprise-AI venture of Abu Dhabi's ATRC (Falcon LLMs, via TII). Chaired by H.E. Faisal Al Bannai. Sovereign, self-hosted AI positioning.

WHAT SHIPS

Design Studio (live public beta: IFC/building-code validation), AiComply (permits), satellite/drone site monitoring, digital twins. Schedule intelligence in development.

WHY THE SEAT STAYS OPEN

Buildings/urban ICP — developers, consultants, municipalities. No contract mechanics, correspondence, notices or claims. UAE-first; no KSA/IKTVA posture.

Threat: medium, rising. If their schedule studio extends into contracts + documents, they become head-on with a sovereign-model advantage — est. 12–24 months away. Speed on the commercial wedge is the counter. Full analysis: `outputs/analyses/competitive-analysis-asset-nova-2026-08-26.md`

A3 · TOUGH QUESTIONS, PREPARED

The three hardest questions

"WHAT IF PEOPLE ARE DELIBERATELY NOT ACTING ON WHAT YOU SURFACE?"

That's the product working. The disposition ("no action — relationship management") becomes a documented decision instead of a silent omission. Right call → it's on record. Wrong call → the trail shows judgment, not negligence. The agent forces awareness, not action. Customer controls register visibility.

"YOUR TWO KEY ASSUMPTIONS ARE BOTH LOW-CONFIDENCE."

The floor case still works: even if the agent catches nothing new, it saves weeks of correspondence reconstruction per event — $5–10M/project, 2–5× ROI. The low-confidence assumptions decide good vs transformational, and the backtest resolves them before any commercial commitment.

"DOCUMENT CRUNCH / TRIMBLE ALREADY DOES CONTRACT AI."

They read the contract. We read the contract **against the live project record** — continuously. Clause highlighter vs commercial early-warning system. Trimble's acquisition validates the market; the dynamic, cross-system piece is unoccupied.